

4-3: User-User Variations and Tuning

Learning Objectives

- Know the main implementation decisions to make in user-user CF
- Understand different options and why they might be selected
- Have a best-practice starting point for configuring a user-user CF recommender

Introduction

- Previous lectures
 - User-user collaborative filtering
 - How user-user CF works
- This lecture
 - Customizations and design decisions

Overview

- Selecting Neighborhoods
- Scoring Items from Neighborhoods
- Normalizing Data
- Computing Similarities
 - Algorithms
 - Tweaks
- Additional Options

Overview

- **Selecting Neighborhoods**
- Scoring Items from Neighborhoods
- Normalizing Data
- Computing Similarities
 - Algorithms
 - Tweaks
- Additional Options

Introduction to Recommender Systems

Selecting Neighborhoods

- All the neighbors
- Threshold similarity or distance
- Random neighbors
- Top- N neighbors by similarity or distance

Introduction to Recommender Systems

How Many Neighbors?

- In theory, the more the better
 - If you have a good similarity metric
- In practice, noise from dissimilar neighbors decreases usefulness
- Between 25 and 100 is often used
 - 30–50 often good for movies
- Fewer neighbors → lower coverage

Introduction to Recommender Systems

Overview

- Selecting Neighborhoods
- **Scoring Items from Neighborhoods**
- Normalizing Data
- Computing Similarities
 - Algorithms
 - Tweaks
- Additional Options

Introduction to Recommender Systems

Scoring Items from Neighborhoods Overview

- Average
- Weighted average
- Multiple linear regression

Weighted average is common, simple, and works well

Introduction to Recommender Systems

- Selecting Neighborhoods
- Scoring Items from Neighborhoods
- **Normalizing Data**
- Computing Similarities
 - Algorithms
 - Tweaks
- Additional Options

Introduction to Recommender Systems

What's wrong with data?

- Users rate differently
- Some rate high, others low
- Some use more of the scale than others
- Averaging ignores these differences
- Normalization compensates for them

Introduction to Recommender Systems

Mean-centering

- Subtract user mean prior to computing
- Re-add when needed

Introduction to Recommender Systems

z-score normalization

- Mean-center, and divide by standard deviation
- Normalizes for the spread across the scale
- Small additional gain in prediction accuracy over mean-centering

Introduction to Recommender Systems

Overview

- Selecting Neighborhoods
- Scoring Items from Neighborhoods
- Normalizing Data
- **Computing Similarities**
 - Algorithms
 - Tweaks
- Additional Options

Introduction to Recommender Systems

Other normalizations

- Subtract item mean
- Subtract item-user mean

Introduction to Recommender Systems

Computing Similarities

- Last time: Pearson correlation

$$s(a, u) = \frac{\sum (r_{ai} - \mu_a)(r_{ui} - \mu_u)}{\sqrt{\sum (r_{ai} - \mu_a)^2} \sqrt{\sum (r_{ui} - \mu_u)^2}}$$

- Usually only over ratings in common
- User normalization not needed
- Spearman rank correlation is Pearson applied to *ranks*
 - Hasn't been found to work as well

Introduction to Recommender Systems

Problem: what about little data?

- Suppose users have 1 rating in common
- Pearson correlation is 1
- Are the users really similar?

Introduction to Recommender Systems

Solution: significance weighting

- Weight similarity by confidence
- Simple approach: multiply by $1/\min(n, 50)$
 - < 50 common ratings: scaled down by # of common ratings
 - ≥ 50 common ratings: unscaled
- Can also do Bayesian damping

Introduction to Recommender Systems

Vector Similarity

- Compute cosine of user vectors in rating space

$$\text{sim}(a, u) = \frac{\mathbf{a} \cdot \mathbf{u}}{|\mathbf{a}| |\mathbf{u}|} = \frac{\sum r_{ai} r_{ui}}{\sqrt{\sum r_{ai}^2} \sqrt{\sum r_{ui}^2}}$$

- With user-mean norm: Pearson correlation!

Introduction to Recommender Systems

Self-weighting Similarity

- Cosine has built-in significance weighting
- Weights proportionally to ratio of common ratings & total ratings (roughly)
- Similar effect as using overall σ_u instead of just over common ratings in Pearson

Introduction to Recommender Systems

Overview

- Selecting Neighborhoods
- Scoring Items from Neighborhoods
- Normalizing Data
- Computing Similarities
 - Algorithms
 - Tweaks
- **Additional Options**

Introduction to Recommender Systems

Pre-computation

- Expensive
- Users move as their ratings change

Introduction to Recommender Systems

Clustering

- Cluster users
- Pick user's cluster to generate predictions
- Doesn't work particularly well

Introduction to Recommender Systems

Baseline Configuration

- Top N neighbors (~ 30)
- Weighted averaging
- User-mean or z-score normalization
- Vector similarity over normalized ratings

Introduction to Recommender Systems

Conclusion

- There are a variety of configuration points
- Current research has suggested some that work well
- Next module will discuss evaluation methods you can use to find good options for your application

4-3: User-User Variations and Tuning

$$\begin{aligned} \text{sim}(a, u) &= \frac{\vec{a} \cdot \vec{u}}{\|\vec{a}\| \|\vec{u}\|} \\ &= \frac{\sum_i \hat{r}_{ai} \cdot \hat{r}_{ui}}{\sqrt{\sum_i \hat{r}_{ai}^2} \sqrt{\sum_i \hat{r}_{ui}^2}} \quad \hat{r}_{ai} = r_{ai} - \mu_a \\ &= \frac{\sum (r_{ai} - \mu_a)(r_{ui} - \mu_u)}{\sqrt{\sum (r_{ai} - \mu_a)^2} \sqrt{\sum (r_{ui} - \mu_u)^2}} \\ &\stackrel{\text{scale by}}{\sim} \frac{|R_a \cap R_u|}{|R_a| |R_u|} \end{aligned}$$

$$\begin{aligned} \text{sim}(a, u) &= \frac{\sum^x (r_{ai} - \mu_a) \sum^y (r_{ui} - \mu_u)}{\sqrt{\sum (r_{ai} - \mu_a)^2} \sqrt{\sum (r_{ui} - \mu_u)^2} + k} \\ &\quad \frac{\sum^x \sum^y}{\sqrt{x^2} \sqrt{y^2}} = 1 \\ &\text{sim}(a, u) \cdot \frac{1}{\min(50, |R_a \cap R_u|)} \end{aligned}$$

$$\mu + \hat{\mu}_i + \hat{\mu}_u$$

$$\hat{\mu}_i = \frac{\sum_u (r_{ui} - \mu)}{\text{\# of ratings}(i)}$$

$$\hat{\mu}_u = \frac{\sum_i (r_{ui} - \mu - \hat{\mu}_i)}{\text{\#}(u)}$$

$$s(a,i) = \frac{\sum_{u \in \mathcal{N}} r_{ui} \cdot \text{sim}(a,u)}{\sum_u \text{sim}(a,u)} + \mu_a$$

$$z_{ui} = \frac{r_{ui} - \mu_u}{\sigma_u}$$

$$s(a,i) = \frac{\sum_u z_{ui} \cdot \text{sim}(a,u)}{\sum_u |\text{sim}(a,u)|} \cdot \sigma_a + \mu_a$$

$$\frac{\sum (r_{ai} - \mu_{ai})(r_{ui} - \mu_u)}{\sqrt{\sum (r_{ai} - \mu_{ai})^2} \sqrt{\sum (r_{ui} - \mu_u)^2}}$$

$$\frac{\sum (r_{ai} - \mu_a)(r_{ui} - \mu_u)}{\sqrt{\sum (r_{ai} - \mu_a)^2} \sqrt{\sum (r_{ui} - \mu_u)^2}}$$