

5-2: Basic Accuracy Metrics

Introduction to Recommender Systems

A little intuition

- Error metrics are usually computed using a “leave one out” methodology
 - Cover up a rating, and try to predict it
- Warning: sometimes this is hard, and evaluators take short cuts (e.g., leave 10% out).

Introduction to Recommender Systems

Goals for Today

- To understand how to compute
 - MAE -- Mean absolute error
 - MSE -- Mean squared error
 - RMSE -- Root mean squared error
- To understand variations on how these may be computed
- To understand where accuracy metrics are useful in general, and the relative merits of each of these three

Introduction to Recommender Systems

Mean Absolute Error (MAE)

- What is error?
 - Divergence of prediction from actual opinion (rating)
 - $P - R$
- Absolute error removes direction
 - $|P - R|$
 - Why? Because two wrongs don't make a right!
- $MAE = \text{Average} (|P - R|)$

$$= \frac{\sum_{ratings} |P - R|}{\# ratings}$$

Introduction to Recommender Systems

Mean Squared Error (MSE)

- Why Squared Error?
 - Removes sign – avoids need for absolute value
 - Penalizes large errors more than small

- $$\frac{\sum ratings (P-R)^2}{\# ratings}$$

- One disadvantage – squared error is not on an intuitive scale ...

Root Mean Squared Error (RMSE)

- $$\sqrt{\frac{\sum ratings (P-R)^2}{\# ratings}}$$

Hold on a moment ...

- We glossed over the summation
 - Usual model – average over all ratings
 - Alternative model – average over user averages
- What's the difference
 - What if one user has 3000 ratings and another 10?
- Advice – consider looking at both – understand what you're comparing to

Comparing Different Algorithms

- What to do when computing MAE in different cases:
 - Remember, must be same data set/scale
 - If coverage is different (different set of user/item pairs for which predictions are available, two choices):
 - Check against common subset
 - Supplement algorithm with default for full coverage

Reflections ...

- In general, all the error metrics move together (good replacements for each other)
- Squared may matter for large scales with some algorithms that have occasional huge errors, but other measures may catch that better
- Benefit – lots of published MAE data for public datasets
- Drawback – error can be dominated by irrelevant parts of the item space

Introduction to Recommender Systems

Looking forward ...

- Next, we look at decision-support metrics

Introduction to Recommender Systems

5-2: Basic Accuracy Metrics