

6-2: Item-Item Algorithm

Structure of Item-Item CF

- Pre-compute item similarities over all pairs of items
- Look for items similar to those the user likes
 - Or has purchased
 - Or has in their basket

Introduction

- We're now into the 2nd major personalized algorithm: item-item CF
- This lecture will discuss the algorithm in more detail
- Also: design space and performance implications

Components

- Item similarity function
- Model builder
- Neighborhood selection strategy
- Item score aggregation function

Item Similarities

- Usually use cosine similarity between item rating vectors
- Often normalize user ratings first
 - Subtract user mean
 - Subtract item mean

$$sim(i, j) = \frac{\sum_{u \in U(i) \cap U(j)} \hat{r}_{ui} \hat{r}_{uj}}{\sqrt{\sum_u \hat{r}_{ui}^2} \sqrt{\sum_u \hat{r}_{uj}^2}}$$

Introduction to Recommender Systems

Scoring Items

- Score is driven by item
- For each item to score:
 - Find similar items the user has rated
 - Compute weighted average of user's ratings

$$p_{ui} = \frac{\sum_{j \in N} sim(i, j) r_{uj}}{\sum_{j \in N} |sim(i, j)|}$$

Introduction to Recommender Systems

Picking Neighbors

- Score formula had a neighborhood N
- Neighbors are usually k most similar items
 - That the user has rated
- Good value of k important
 - k too small \rightarrow inaccurate scores
 - k too large \rightarrow too much noise (low-similarity items)
 - $k=20$ often works well

Introduction to Recommender Systems

Building the Model

- Pre-compute similarities for all pairs of items
 - Item stability makes similarity pre-computation feasible
- Naïvely: $O(|I|)^2$
 - If symmetric: only need to compute one direction
 - Sometimes can skip pairs

Introduction to Recommender Systems

Truncating the Model

- Don't need to keep the whole I^2 model
- But need enough neighbors to find neighbors at score time
 - Since user hasn't rated everything, need $M \gg k$ neighbors per item in model
- Balance memory use with accuracy and coverage
 - Mild runtime impact, if neighbors are sorted

Introduction to Recommender Systems

Tuning the model

- Tune using cross-validation
- Need to find good values for
 - baseline and normalization
 - similarity function (or just use cosine)
 - neighborhood size k
 - model size M
 - sometimes similarity is damped, but often doesn't help much

Introduction to Recommender Systems

Conclusion

- Item-item is efficient and straightforward
- A few parameters need tuning for specific data, domain
- Next: more tweaks
 - applying to unary data (implicit feedback)
 - repurposing and hybridization

Introduction to Recommender Systems

6-2: Item-Item Algorithm

Introduction to Recommender Systems

Introduction to Recommender Systems