

7-3: Training SVDs

Gradient Descent

- Convex optimization method
- We won't cover gradient descent in general here
 - that's multivariate calculus
- Basic idea: how far off we are at one point (prediction) tells us how to improve the prediction
- We'll use *stochastic* gradient descent

Introduction

- SVDs are slow
- SVDs require care dealing with missing data
- This video: gradient descent
 - Much faster
 - Deals well with missing data

Algorithm Structure

- Initialize values to train (item/user feature vectors) to arbitrary starting point
 - Must be non-zero
- Try to predict each rating in data set
- Use error and *update rule* to update values for next rating/sample
- Iterate until convergence
 - Stops moving
 - Iterated enough times

FunkSVD

- Train features one at a time
 - Train first feature until convergence
 - Then train the next
- Ignore missing values (mostly)
- Learn offset from personalized mean
- Fold singular values into left/right matrices

$$R = B + UV^T$$

$$p_{a,i} = b_{a,i} + \sum_f u_{a,f} v_{i,f}$$

Introduction to Recommender Systems

Parameters

$$\Delta v_{if} = \lambda((r_{a,i} - p_{a,i})u_{af} - \gamma v_{if})$$

$$\Delta u_{af} = \lambda((r_{a,i} - p_{a,i})v_{if} - \gamma u_{af})$$

- λ – learning rate (how fast to train)
 - Higher – faster training, lower accuracy
- γ – regularization factor
 - prevents ‘overfitting’ by encouraging small values

Introduction to Recommender Systems

Where'd this come from?

- Gradient descent to minimize (global) RMSE of final model
- Update is derivative of squared error with respect to user or item feature value
 - With regularization term
- General idea: use gradient descent (or other method) to find values that minimize the error (cost function)
- Caveat: no longer a true SVD

Introduction to Recommender Systems

Extending: SVD++

- Minimize-the-error model is very general
 - Gradient descent can solve many of them
- Extension: incorporate has-rated data
- $N(a)$ is items rated (or purchased) by a


$$p_{a,i} = b_{a,i} + (u_a + \sqrt{|N(a)|} \sum_{i \in N(a)} y_i) v_i^T$$

Introduction to Recommender Systems

Conclusion

- Gradient descent allows efficient approximation of the SVD from known data
- Extensible to additional data – minimize the error over some predictor
 - Limited by error's ability to capture what's important to recommendation problem
- Many models exist using similar methods

Introduction to Recommender Systems



7-3: Training SVDs

Introduction to Recommender Systems

$$p_{ai} = b_{ai} + \sum_f u_{af} v_{if} \quad \begin{array}{l} \text{initialize to 0.1} \\ \text{or nonzero, small random} \end{array}$$
$$e_{ai} = r_{ai} - p_{ai}$$
$$\Delta v_{if} = \lambda (e_{ai} u_{af} - r v_{if}) \quad \begin{array}{l} \text{small } (\sim 0.1) \\ \text{regularization} \end{array}$$
$$\Delta u_{af} = \lambda (e_{ai} v_{if} - r u_{af}) \quad \begin{array}{l} \text{learning rate (0.001)} \end{array}$$